

World-space accuracy, 100-frame segments

HOI4D, 157 sequences, detector boxes shared across every row, one scorer. Lower is better.

Method	Regime	C-abs	C-rr	WA-MPJPE	W-MPJPE	seg.
Ours (feedforward)	online, no SLAM	35.4	26.5	27.0 <i>was 34.2</i>	70.7 <i>was 113.2</i>	471
WiLoR + SLAM	offline + SLAM	83.4	27.2	35.2	129.0	468
HaWoR	offline + SLAM	87.7	32.6	40.3	133.6	471
HaMeR + SLAM	offline + SLAM	87.9	30.3	36.4	135.8	468
HaPTIC + SLAM	offline + SLAM	(157.1)	29.6	36.5	(138.1)	468
Dyn-HaMR	offline optim.	(1336.7)	59.9	49.0	(195.4)	468

What changed, and the control

The predicted camera trajectory was being applied backwards: the tensor is world-to-camera, the code that emits it labels it camera-to-world. Over 1570 clips the rotation error is 2.1° as emitted against 0.36° inverted, and the best-fit translation scale is -0.944 . Correcting it moves W-MPJPE 113.2 to 70.7 and WA 34.2 to 27.0.

C-abs and C-rr do not move by a decimal. Neither reads a camera pose, so neither may change, and neither does. No weight was retrained and no baseline cell moves: none of them consumes this tensor.

Parenthesised cells are measured but not attributable: neither method recovers metric depth under input matching.